

Method and measured results

One sampled video latent, decoded natively and with 128 / 384 px of horizontal wrap context.



Full ERP, frame 97. The wrap is at the left/right edges. Perspective views on the following pages centre that boundary.

Wrap -> decode -> crop

Append columns from the opposite side of the latent to each horizontal edge. Decode the wider latent with the native H3 VAE, then crop the added margins. At a spatial scale of 16, 128 / 384 px corresponds to 8 / 24 latent columns per side.

Sampling and model weights are unchanged; no additional diffusion pass. The wider canvas also changes the native decoder tile layout, so this comparison includes both effects.

Matched settings

1536 x 672 | 243 frames | 24 fps | 50 steps | reviewed 360 LoRA at 1.0 | H3 BF16 | VAE FP16

Measurement	Native	Circular 128	Circular 384	Change vs native
Temple: boundary MAE	5.841	4.098	4.206	-29.8% / -28.0%
Forest: boundary MAE	12.212	8.375	8.591	-31.4% / -29.7%
Temple: decode time (H200)	13.29 s	14.75 s	19.57 s	+1.46 s / +6.28 s

How improvement is measured


Boundary MAE is the mean absolute RGB difference between the first and last image columns, averaged over rows, channels and all 243 decoded frames (0-255 scale). Lower means less pixel mismatch at the wrap; it is not a percentage of visual improvement.

Visual checks use matched frames, spherical views and playback. Fine geometry and polar distortion remain; two scenes do not establish general reliability. Temporal pixel change is not motion-compensated and does not establish reduced flicker. Refinement/upscaling is still untested.


The seam at different viewing angles

Native, circular 128 and circular 384. Same frame (97), yaw (180 degrees) and horizontal field of view (75 degrees).


Eye level / pitch 0 degrees



Looking up / pitch +45 degrees



Looking down / pitch -45 degrees



Seam detail / circular 128

Native and circular 128 at identical magnification. Same frame (97), yaw 180 degrees, pitch 0 degrees, 20-degree field of view.



Close-up comparison

The vertical boundary runs through the carved face in the centre of each panel. Wrapping changes how the facial features and surrounding stone join across that boundary; residual shape differences remain.

Both sides use the same projection and resampling of the saved PNG frames. This is magnification of existing detail, without sharpening or generated enhancement.

Seam detail / circular 384

Native and circular 384 at identical magnification. Same frame (97), yaw 180 degrees, pitch 0 degrees, 20-degree field of view.



Close-up comparison

The vertical boundary runs through the carved face in the centre of each panel. Wrapping changes how the facial features and surrounding stone join across that boundary; residual shape differences remain.

Both sides use the same projection and resampling of the saved PNG frames. This is magnification of existing detail, without sharpening or generated enhancement.